

# GradientDescent

November 26, 2024

## 1 Data Generation

```
[1]: %matplotlib inline
import numpy as np
import numpy.linalg as la
import matplotlib.pyplot as plt

dim_theta = 10
data_num = 1000
scale = .1

theta_true = np.ones((dim_theta,1))
print('True theta:', theta_true.reshape(-1))

A = np.random.uniform(low=-1.0, high=1.0, size=(data_num,dim_theta))
y_data = A @ theta_true + np.random.normal(loc=0.0, scale=scale,
↪size=(data_num, 1))

A_test = np.random.uniform(low=-1.0, high=1.0, size=(50, dim_theta))
y_test = A_test @ theta_true + np.random.normal(loc=0.0, scale=scale, size=(50,
↪1))
```

True theta: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

## 2 Solving for the exact mean squared loss (solving $Ax = b$ )

```
[85]: '''
Using numpy's lstsq() to get the closed-form least squares solution
'''
theta_pred = la.lstsq(A, y_data)[0]

print('Empirical theta: ', theta_pred.reshape(-1))
```

Empirical theta: [1.01117113 0.99879624 0.99707909 1.01475117 1.00365197  
0.98530922  
1.00109998 0.99715709 0.99246624 0.99988042]

### 3 SGD Variants Noisy Function

```
[11]: batch_size = 1
      max_iter = 1000
      lr = 0.001
      theta_init = np.random.random((10,1)) * 0.1

[15]: def noisy_val_grad(theta_hat, data_, label_, deg_=2.):
      gradient = np.zeros_like(theta_hat)
      loss = 0

      for i in range(data_.shape[0]):
          x_ = data_[i, :].reshape(-1,1)
          y_ = label_[i, 0]
          err = np.sum(x_ * theta_hat) - y_

          grad = deg_ * np.abs(err) ** (deg_ - 1) * np.sign(err) * x_
          l = np.abs(err) ** deg_

          loss += l / data_.shape[0]
          gradient += grad / data_.shape[0]

      return loss, gradient
```

### 4 Running SGD Variants

```
[95]: #@title Parameters
      deg_ = 2 #@param {type: "number"}
      num_rep = 15 #@param {type: "integer"}
      max_iter = 1000 #@param {type: "integer"}
      fig, ax = plt.subplots(figsize=(6, 6))
      best_vals = {}
      test_exp_interval = 50 #@param {type: "integer"}
      grad_artificial_normal_noise_scale = 0. #@param {type: "number"}

      # dictionary for keeping track of parameters
      params = {}

      for method_idx, method in enumerate(['adam', 'sgd', 'adagrad']):
          test_loss_mat = []
          train_loss_mat = []

          for replicate in range(num_rep):
              if replicate % 20 == 0:
                  print(method, replicate)
```

```

if method == 'adam':
    beta_1 = 0.9
    beta_2 = 0.999
    m = np.zeros((dim_theta, 1))
    v = np.zeros((dim_theta, 1))
    epsilon = 1e-8

if method == 'adagrad':
    epsilon = 1e-8
    squared_sum = np.zeros((dim_theta, 1))

theta_hat = theta_init.copy()
test_loss_list = []
train_loss_list = []

for t in range(max_iter):
    idx = np.random.choice(data_num, batch_size) # Split data
    train_loss, gradient = noisy_val_grad(theta_hat, A[idx,:],
    ↪ y_data[idx:], deg=deg_)
    artificial_grad_noise = np.random.randn(10, 1) *
    ↪ grad_artificial_normal_noise_scale + np.sign(np.random.random((10, 1)) - 0.
    ↪ 5) * 0.
    gradient = gradient + artificial_grad_noise
    train_loss_list.append(train_loss)

    if t % test_exp_interval == 0:
        test_loss, _ = noisy_val_grad(theta_hat, A_test[:, :], y_test[:, :
    ↪ ], deg=deg_)
        test_loss_list.append(test_loss)

    if method == 'adam':
        m = beta_1 * m + (1-beta_1) * gradient
        v = beta_2 * v + (1-beta_2) * (gradient * gradient)
        m_hat = m / (1 - beta_1 ** (t+1))
        v_hat = v / (1 - beta_2 ** (t+1))
        theta_hat = theta_hat - lr * (m_hat/(np.sqrt(v_hat)+epsilon))

    elif method == 'adagrad':
        squared_sum = squared_sum + (gradient * gradient)
        theta_hat = theta_hat - lr * gradient * (1/np.
    ↪ sqrt(squared_sum+epsilon))

    elif method == 'sgd':
        theta_hat = theta_hat - lr * gradient

test_loss_mat.append(test_loss_list)
train_loss_mat.append(train_loss_list)

```

```

print(method, 'done')
x_axis = np.arange(max_iter)[::test_exp_interval]

test_loss_np = np.array(test_loss_mat)

'''
Hints:
1. Use test_loss_np in np.mean() with axis = 0
'''
test_loss_mean = np.mean(test_loss_np, axis=0)

'''
Hints:
1. Use test_loss_np in np.std() with axis = 0
2. Divide by np.sqrt() using num_rep as a parameter
'''
test_loss_se = np.std(test_loss_np, axis=0) / np.sqrt(num_rep)

plt.errorbar(x_axis, test_loss_mean, yerr=2.5*test_loss_se, label=method)
best_vals[method] = min(test_loss_mean)

# this only gets the parameters for one replicate (the last one) but it's
→okay because we only need one to answer q3.
params[method] = theta_hat

best_vals = { k: int(v * 1000) / 1000. for k,v in best_vals.items() } # A weird
→way to round numbers
plt.title(f'Test Loss \n(objective degree: {deg_}, best values: {best_vals})')
plt.ylabel('Test Loss')
plt.legend()
plt.xlabel('Updates')

```


```

adam 0
adam done
sgd 0
sgd done
adagrad 0
adagrad done

```

```
[95]: Text(0.5, 0, 'Updates')
```

Test Loss  
(objective degree: 2, best values: {'adam': 1.941, 'sgd': 1.008, 'adagrad': 3.577})



```
[94]: print(params["adam"].reshape(-1))
```

```
[0.35388391 0.27269951 0.31367533 0.36023661 0.28976016 0.36528544  
0.33467022 0.36753296 0.29897261 0.30564255]
```

```
[ ]:
```